

RL

Reuben Brasher

December 6, 2022

Pictures of games

k -armed Bandits

Reinforcement Learning

k-armed Bandit



Figure: From https://en.wikipedia.org/wiki/Slot_machine

Markov Decision Process




Figure: From https://en.wikipedia.org/wiki/Reinforcement_learning

Problem

Each round t an agent may choose an action from k possible. The agent receives a reward sampled from a distribution conditioned on the action.

$$P(R_t|A_t)$$

The objective of the game is to learn which action will give the highest expected reward.

q_*

- ▶ If only we knew the *value* of each action

$$q_*(a) = \mathbb{E}[R_t | A_t = a].$$

q_*

- ▶ If only we knew the *value* of each action

$$q_*(a) = \mathbb{E}[R_t | A_t = a].$$

- ▶ We do not. We know $Q_t(a)$.

Explore vs. Exploit

- ▶ Greedy strategy is always choose current best

$$\operatorname{argmax}_a Q_t(a)$$

Explore vs. Exploit

- ▶ Greedy strategy is always choose current best

$$\operatorname{argmax}_a Q_t(a)$$

- ▶ ε -Greedy strategy is to choose uniformly randomly probability ε , and to follow greedy strategy otherwise.

Explore vs. Exploit

- ▶ Greedy strategy is always choose current best

$$\operatorname{argmax}_a Q_t(a)$$

- ▶ ε -Greedy strategy is to choose uniformly randomly probability ε , and to follow greedy strategy otherwise.
- ▶ Upper confidence bound strategy is to choose

$$\operatorname{argmax}_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

Finite Markov Decision Process

Agent and environment interact to produce a trajectory

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, \dots$$

State and reward depend on previous state and agent action

$$p(s', r|s, a) = \Pr(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$$

G_t

► *Return*

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T$$

G_t

► *Return*

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T$$

► *Discounted return*

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Policy, value and action value function

Agent responds to environment by sampling from *policy* $\pi(a|s)$.
Value of state s is

$$v_{\pi}(s) = \mathbb{E}[G_t | S_t = s]$$

Action value function is

$$q_{\pi}(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$$

If only we knew...

The *optimal policy* π_* or
the *optimal value* of state s

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

or the *optimal action value function*

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

Note that

$$q_*(s, a) = \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a]$$

Bellman optimality equations

- ▶ Bellman equation is

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

Bellman optimality equations

- ▶ Bellman equation is

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s',r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

- ▶ Bellman optimality equation for value is

$$v_*(s) = \max_a \sum_{s',r} p(s', r|s, a) [r + \gamma v_*(s')]$$

Bellman optimality equations

- ▶ Bellman equation is

$$v_{\pi}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]$$

- ▶ Bellman optimality equation for value is

$$v_*(s) = \max_a \sum_{s', r} p(s', r|s, a) [r + \gamma v_*(s')]$$

- ▶ Bellman optimality equation for action value is

$$q_*(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma \max_{a'} q_*(s', a')]$$

Finding optimal policy

- By policy iteration (pg 121)

$$\pi_0 \rightarrow V_{\pi_0} \rightarrow \pi_1 \rightarrow V_{\pi_1} \rightarrow \dots \rightarrow \pi_* \rightarrow V_*$$

Finding optimal policy

- By policy iteration (pg 121)

$$\pi_0 \rightarrow V_{\pi_0} \rightarrow \pi_1 \rightarrow V_{\pi_1} \rightarrow \dots \rightarrow \pi_* \rightarrow V_*$$

- By computing q_*

Finding optimal policy

- ▶ By policy iteration (pg 121)

$$\pi_0 \rightarrow v_{\pi_0} \rightarrow \pi_1 \rightarrow v_{\pi_1} \rightarrow \dots \rightarrow \pi_* \rightarrow v_*$$

- ▶ By computing q_*
- ▶ By computing v_*

Temporal Difference

Iteratively update V to estimate v_π .

Update V by

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

Why wait though?

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

SARSA

Iteratively update Q to estimate q_* . (pg 151)

Update Q by

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Q-learning

Iteratively update Q to estimate q_* . (pg 153)

Update Q by

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

Policy Gradient

Parameterize policy π by θ , for example

$$\pi(a|s, \theta) = \frac{e^{h(s,a,\theta)}}{\sum_b e^{h(s,b,\theta)}}$$

where $h(s, a, \theta)$ is ANN or similar.

Then π has a gradient with respect to θ , and v_π can be improved by gradient ascent.

Actor critic

Parameterize policy π by θ , and estimate v by w

$$\pi(a|s, \theta) = \frac{e^{h(s,a,\theta)}}{\sum_b e^{h(s,b,\theta)}}$$

where $h(s, a, \theta)$ is ANN or and so is $v(s, w)$.

Then π has a gradient with respect to θ , and v with respect to w .
Improve v and π in turn by applying gradient updates.

Futher reading

Sutton, Barto, Bach, et al., 1998





Narasimhan, Kulkarni, and Barzilay, 2015

Li et al., 2016

Mnih et al., 2015

Zhu et al., 2016

References I

-  Li, Jiwei et al. (2016). “Deep reinforcement learning for dialogue generation”. In: *arXiv preprint arXiv:1606.01541*.
-  Mnih, Volodymyr et al. (2015). “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540, pp. 529–533.
-  Narasimhan, Karthik, Tejas Kulkarni, and Regina Barzilay (2015). “Language understanding for text-based games using deep reinforcement learning”. In: *arXiv preprint arXiv:1506.08941*.
-  Sutton, Richard S, Andrew G Barto, Francis Bach, et al. (1998). *Reinforcement learning: An introduction*. MIT press.

References II



Zhu, Yuke et al. (2016). “Target-driven visual navigation in indoor scenes using deep reinforcement learning”. In: *arXiv preprint arXiv:1609.05143*.